

L04 Sockets

Definition

- The point where a local application process attaches to the network
- Interface Between network and application
- Created by application

Socket Working Process

- Being created
- Being connected to the network
- Changing messages
- Being closed

Ports

Used to identify different applications(processes) in the same host

Socket Implementation

1. Socket Family

- PF_INET
 - **INET** denotes the internet family. If u want to communicate with a remote host, u should use this.
 - An **INET** socket is bound to a specific IP address.
 - INET sockets sit at the top of a full TCP/IP stack, with traffic congestion algorithms, backoffs and the like to handle
- PF_UNIX
 - An **UNIX** socket is bound to a special file on our file system.
 - We use it as an a lightweight alternative to an **INET** socket via loopback, when u need communication between processes on the same host.
 - everything is designed to be local to the machine, so its code is much simpler and the communication is faster
- PF_PACKET
 - Used if u want to play with packets at the protocol level, if u are implementing ur own protocol
 - denotes direct access to the network interface (i.e., it bypasses the TCP/IP protocol stack)

Remark AF_ and PF_ have the same values. And u can just use them alternatively, but AF_ seem to be more prevalent.

2. Socket Type

- SOCK_STREAM: denotes a byte stream
- SOCK_DGRAM: an alternative that denotes a message oriented service, such as that provided by UDP

Remark PF_INET and SOCK_STREAM imply **TCP**

3. Socket in Code

Generic

IP Specific

```
• struct sockaddr
{
    unsigned short sa_family; /* Address family (e.g., AF_INET) */
    char sa_data[14];        /* Protocol-specific address information */
};

• struct sockaddr_in
{
    unsigned short sin_family; /* Internet protocol (AF_INET) */
    unsigned short sin_port;   /* Port (16-bits) */
    struct in_addr sin_addr;    /* Internet address (32-bits) */
    char sin_zero[8];          /* Not used */
};

struct in_addr
{
    unsigned long s_addr;      /* Internet address (32-bits) */
};
```

sockaddr	Family	Blob		
	2 bytes	2 bytes	4 bytes	8 bytes
sockaddr_in	Family	Port	Internet address	Not used

4. Helper Functions

- Printable String to IP address

int inet_pton()

- Binary to printable string

const char * inet_ntop

- convert the integer from host byte order to network byte order

uint16_t htons (uint16_t hostshort)

- convert the integer from network byte order to host byte order

uint16_t ntohs (uint16_t netshort)

- connect socket

int connect(int socket, const struct sockaddr *address, socklen_t address_len)

socket is the descriptor

- Sending data

ssize_t send(socket,buf,len,flags)

- return the number of bytes being copied to the OS kernel for transmission
- blocking semantics of send

- Receiving Responses

ssize_t recv(int sockfd , void buf , size_t len , int flags);

- Receives at least one byte
- 0 is reserved when the other side closed the socket
- Does not receive the same number of bytes that were sent via send
- You have to keep reading from the socket until you've received all the bytes you need
- returns the length of the message or datagram in bytes

- Bind (Only for server)

int bind(int sockfd , const struct sockaddr my_addr , socklen_t addrlen)

my_addr

- struct sockaddr_in for IPv4
- cast (struct sockaddr_in *) to (struct sockaddr)

Remark Binding a socket means assigning an address and port number to the socket

- Listen(Only for server)

int listen(int sockfd , int backlog)

backlog Number of pending connections to queue

- Accept(Only for server)

int accept(int sockfd , struct sockaddr addr , socklen_t addrlen)

Eg: int isock = sockfd , (struct sockaddr_in *) caddr , clen);

Socket Options

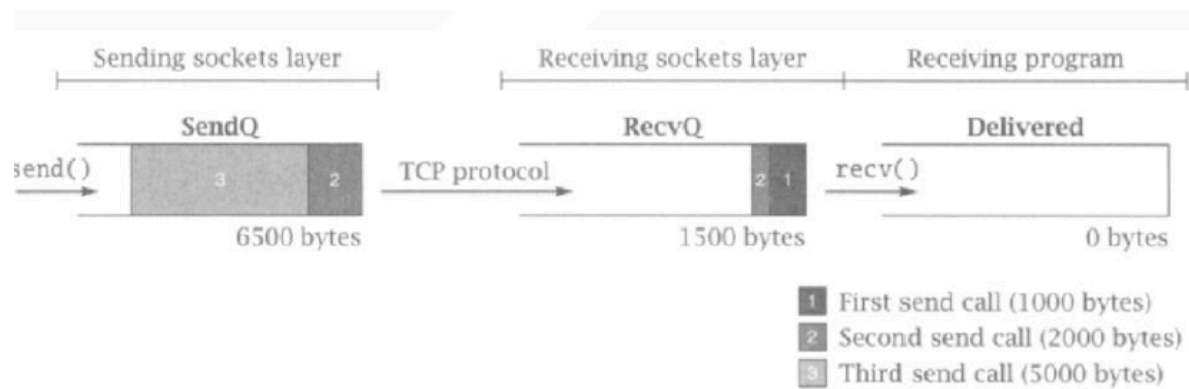
GETSOCKOPT(2)	Linux Programmer's Manual	GETSOCKOPT(2)
NAME		
getsockopt, setsockopt – get and set options on sockets		
SYNOPSIS		
<pre>#include <sys/types.h> /* See NOTES */ #include <sys/socket.h> int getsockopt(int sockfd, int level, int optname, void *optval, socklen_t *optlen); int setsockopt(int sockfd, int level, int optname, const void *optval, socklen_t optlen);</pre>		
Level	Option	Description
SOL_SOCKET	SO_SNDBUF	Send buffer size
	SO_REUSEADDR	Allow TCP port to be reused immediately
	SO_RCVTIMEO	Set a recv() timeout
	SO_SNDTIMEO	Set a send() timeout

Socket Internals

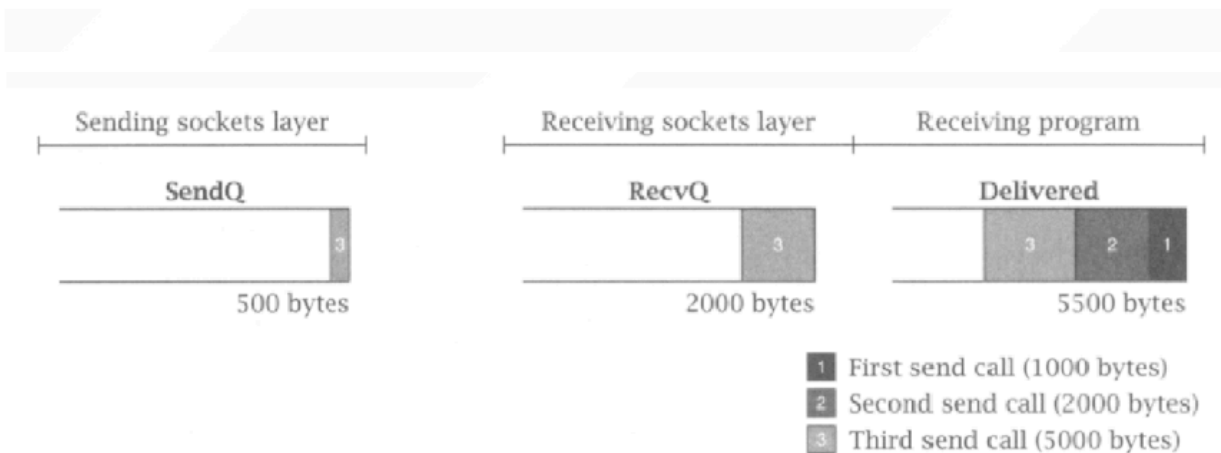

Suppose the following occasion:

1. rv = connect(s, ...);
2. rv = send(s, buffer0, 1000, 0);
3. rv = send(s, buffer1, 2000, 0);
4. rv = send(s, buffer2, 5000, 0);
5. close(s);

At first, we are sending smoothly.



But then something strange happens.



Our sending process seems to be blocked....But why?

Preventing Deadlock!!

Like we just mentioned when we were learning how to send a message with a socket, it is blocked while sending.

WHEN DOES BLOCKING OCCUR?

- SendQ size: **SQS**
- RecvQ size: **RQS**
- `send(s, buffer, n, 0);`
 - $n > \text{SQS}$: blocks until $(n - \text{SQS})$ bytes xfered to RecvQ
 - If $n > (\text{SQS} + \text{RQS})$, blocks until receiver calls `recv()` enough to read in $n - (\text{SQS} + \text{RQS})$ bytes
- How does this lead to deadlock?
 - Trivial cause: both sides call `recv()` w/o sending data

This is a [link](#) I found regardidng this topic. U might found it useful.