

L7 Remote Procedure Call

Why RPC?

Retain the feel of writing centralized code, easy to program network communication that makes client-server communication

Distributed programming is hard ==> abstraction

Goal of RPC

To make communication appear like a local procedure call : transparency for procedure calls

RPC EXAMPLE

Local computing	Remote computing
<pre>X = 3 * 10; print(X) > 30</pre>	<pre>server = connectToServer(S); Try: X = server.mult(3,10); print(X) Except e: print "Error!" > 30 or > Error</pre>

Difficulties

- Heterogeneity

Client needs to communicate with server, but how to solve it when servers are of different types?

1. They may have different data representations.

- represent data using different sizes
- different byte ordering

X86-64: little endian some: big endian

- represent floating numbers differently
- have different data alignment

2. Language support varies

- Many programming language have no inbuilt concept of RPC(C, C++, eariler Java)
- Some have: Python, Haskell, Go

- Failure
 - Message dropped
 - client, server, network fails
- Performance
 - RPC in a datacenter is 1000 times slower than procedure calls

Solution

Write an interface description in the IDL

Use IDL Compiler to convert native data types with machine-independent byte streams

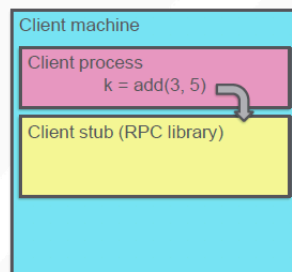
Client stub: forward procedure calls as a request to sever

Server stub: dispatch RPC to its implementation

A Day of an RPC

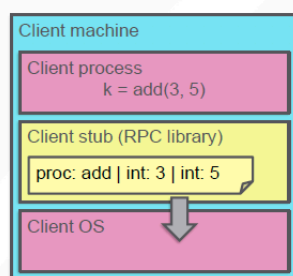
A DAY IN THE LIFE OF AN RPC

1. Client calls stub function (pushes params onto stack)



A DAY IN THE LIFE OF AN RPC


1. Client calls stub function (pushes params onto stack)
2. Stub marshals parameters to a network message



A DAY IN THE LIFE OF AN RPC

2. Stub marshals parameters to a network message

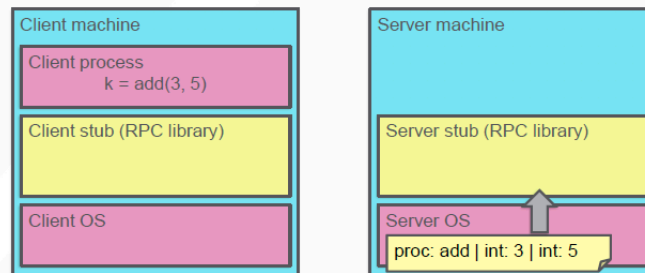
3. OS sends a network message to the server



A DAY IN THE LIFE OF AN RPC

3. OS sends a network message to the server

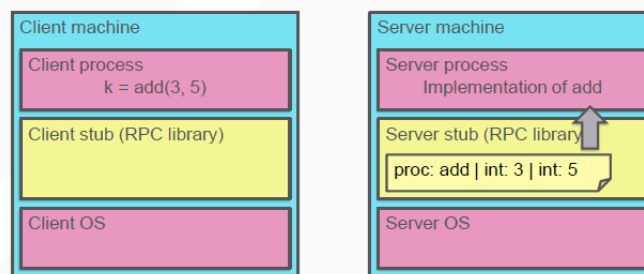
4. Server OS receives message, sends it up to stub



A DAY IN THE LIFE OF AN RPC

4. Server OS receives message, sends it up to stub

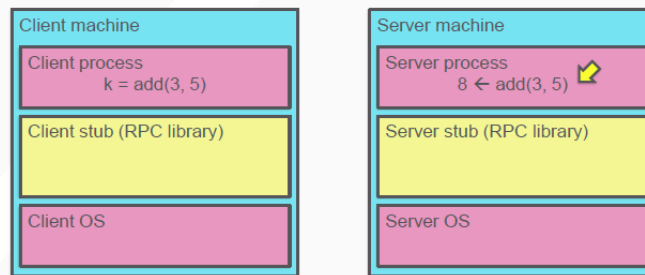
5. Server stub unmarshals params, calls server function



A DAY IN THE LIFE OF AN RPC

5. Server stub unmarshals params, calls server function

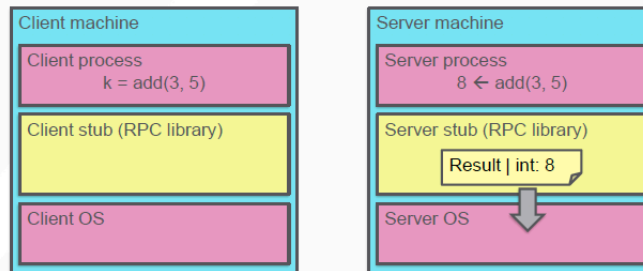
6. Server function runs, returns a value



A DAY IN THE LIFE OF AN RPC

6. Server function runs, returns a value

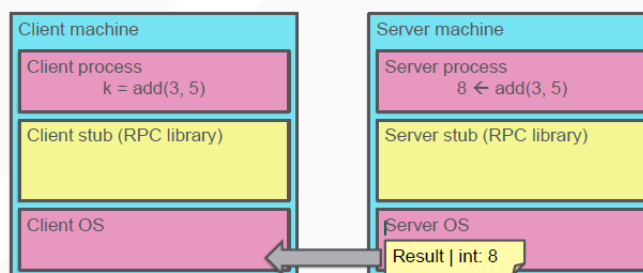
7. Server stub marshals the return value, sends msg



A DAY IN THE LIFE OF AN RPC

7. Server stub marshals the return value, sends msg

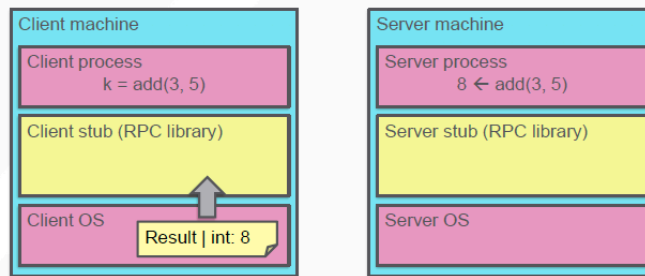
8. Server OS sends the reply back across the network



A DAY IN THE LIFE OF AN RPC

8. Server OS sends the reply back across the network

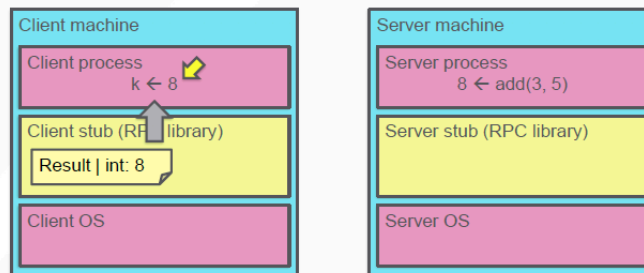
9. Client OS receives the reply and passes up to stub



A DAY IN THE LIFE OF AN RPC

9. Client OS receives the reply and passes up to stub

10. Client stub unmarshals return value, returns to client



Server Stub

- Dispatcher
 - Receives a client's RPC request, identifies appropriate server-side method to invoke
- Skeleton
 - Unmarshals parameters to server-native types
 - Calls the local server procedure
 - Marshals the response, sends it back to the dispatcher

Handling failure

Four kinds of failures may occur

- Client crash and reboot
- Packet dropped (packet loss, broken routing)
- Server crash and reboot
- network or server by very slow

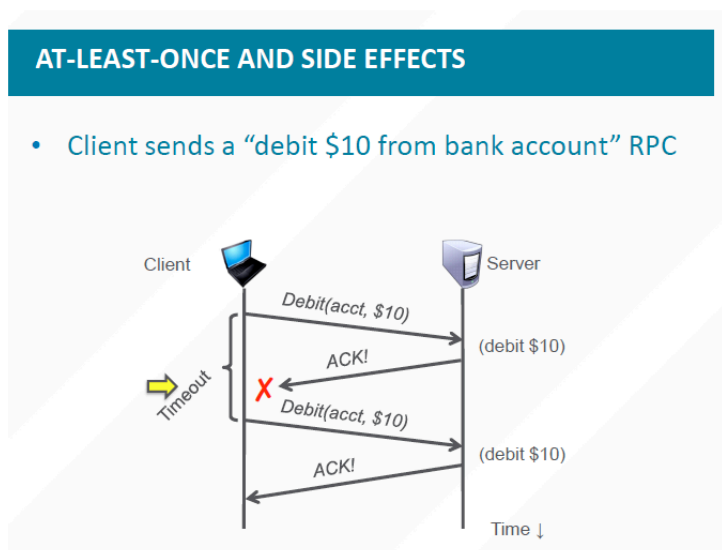
The latter three look the same to the client

It is hidden from the client.

Solution1: AT LEAST ONCE SCHEME

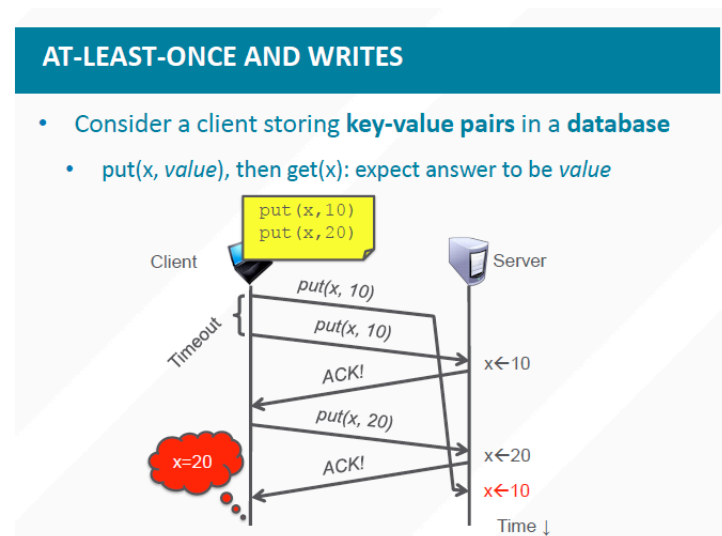
simplest idea for handling failures

1. Client stubs wait for a response (An acknowledge message from the server stub)
2. If no response arrives after a fixed timeout, the client stub resends the request.
3. If repeat the request for several times and still no response, return an error to the application



AT_LEAST_ONCE scheme is okay if there are only read-only operations with no side effects.

If the application has its own functionality to cope with duplication and reordering.



Solution2: AT-MOST-ONCE

Idea: Server RPC code detects duplicate requests

Returns previous reply instead of re running handler

How to detect a duplicate request?

Server sees same function, same arguments twice ==> Is it a duplicate request?

No!!! Some applications legitimately submit the same function with same arguments, twice in a row.

Solution

Client includes a unique transaction ID(xid) with each of its RPC requests

Uses same xid for retransmitted requests

- *How to detect a duplicate request?*
- Client includes unique **transaction ID (xid)** with each one of its RPC requests
- Client uses **same xid** for retransmitted requests

```
At-Most-Once Server  
if seen[xid]:  
    retval = old[xid]  
else:  
    retval = handler()  
    old[xid] = retval  
    seen[xid] = true  
return retval
```

How to ensure xid unique?

1. Combine a unique client ID(IP address) with the current time of the day
2. Combine unique client ID with a sequence number
3. Big random number

Problem with AT-MOST-ONCE

seen and old arrays will grow without bound

Observation

By construction, when the client gets a response to a particular xid , it will never re-send it

Client should tell server "I'm done with xid x --delete it"

Have to tell server each and every retired xid -- paggyback

Significant overhead if many RPCs are in flight

Solution

Client includes “seen all replies $\leq X$ ” with every RPC

Each one of these is cumulative: later seen messages subsume earlier ones

Problem

How to handle a duplicate request while the original is still executing?

Idea

Add a pending flag per executing RPC Server waits for the procedure to finish, or ignores

Problem

Server may crash and restart

Solution

Write seen[], old[] to disk

Summary

RPC SEMANTICS

Delivery Guarantees			RPC Call Semantics
Retry Request	Duplicate Filtering	Retransmit Response	
No	NA	NA	<i>Maybe</i>
Yes	No	Re-execute Procedure	<i>At-least once</i>
Yes	Yes	Retransmit reply	<i>At-most once</i>

- RPC everywhere!
- Necessary issues surrounding machine heterogeneity
- Subtle issues around handling failures